

# Do not distribute course material

You may not and may not allow others to reproduce or distribute lecture notes and course materials publicly whether or not a fee is charged.

<https://mpatacchiola.github.io/blog/2020/07/31/gaussian-mixture-models.html>

Chapter 9 in Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow

<http://cs229.stanford.edu/notes2020spring/cs229-notes7a.pdf>

<http://cs229.stanford.edu/notes2020spring/cs229-notes7b.pdf>

pages 179-181 in <http://ciml.info/>

# Clustering, K-Means and EM

---


INTRODUCTION TO MACHINE LEARNING

PROF. LINDA SELLIE

THANKS TO PROF RANGAN FOR SOME OF THE SLIDES

# Outline

---



Motivating Examples: Document clustering, image segmentation, image compression



- ❑ K-means
- ❑ K++-means (how to initialize the parameters before starting the algorithm)
- ❑ Hyperparameter K
- ❑ (On our own) K-means for document clustering

# Unsupervised Machine Learning

$$\{(\mathbf{x}^{(1)}, \cancel{y}^{(1)}), (\mathbf{x}^{(2)}, \cancel{y}^{(2)}), \dots, (\mathbf{x}^{(N)}, \cancel{y}^{(N)})\}$$

$$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$$

$$\mathbf{x}^{(i)} \in \mathbb{R}^D$$





Pair share: how many clusters should we make?

# Unsupervised Machine Learning

$$\{(\mathbf{x}^{(1)}, \cancel{y}^{(1)}), (\mathbf{x}^{(2)}, \cancel{y}^{(2)}), \dots, (\mathbf{x}^{(N)}, \cancel{y}^{(N)})\}$$

$$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$$

$$\mathbf{x}^{(i)} \in \mathbb{R}^D$$





Pair share: how many clusters should we make?

# Unsupervised Machine Learning

$$\{(\mathbf{x}^{(1)}, \cancel{y}^{(1)}), (\mathbf{x}^{(2)}, \cancel{y}^{(2)}), \dots, (\mathbf{x}^{(N)}, \cancel{y}^{(N)})\}$$

$$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$$

$$\mathbf{x}^{(i)} \in \mathbb{R}^D$$



Pair share: how many clusters should we make?

The goal is to have examples in the same cluster be “close” to each other

# Some clustering applications

---

- Customer segmentation based on their purchases and activities. Allows targeted marketing for different clusters
- Dimensionality reduction: If there are  $k$  cluster each example will have  $k$  new features. Each feature is a measure of how well the example fits into a cluster
- Impute missing values
- Anomaly detection (aka outlier detection)
- Semi-supervised learning (you receive a small amount of labeled data). Label the unlabeled data in the cluster according to the labeled data
- Search engines
- Segmentation



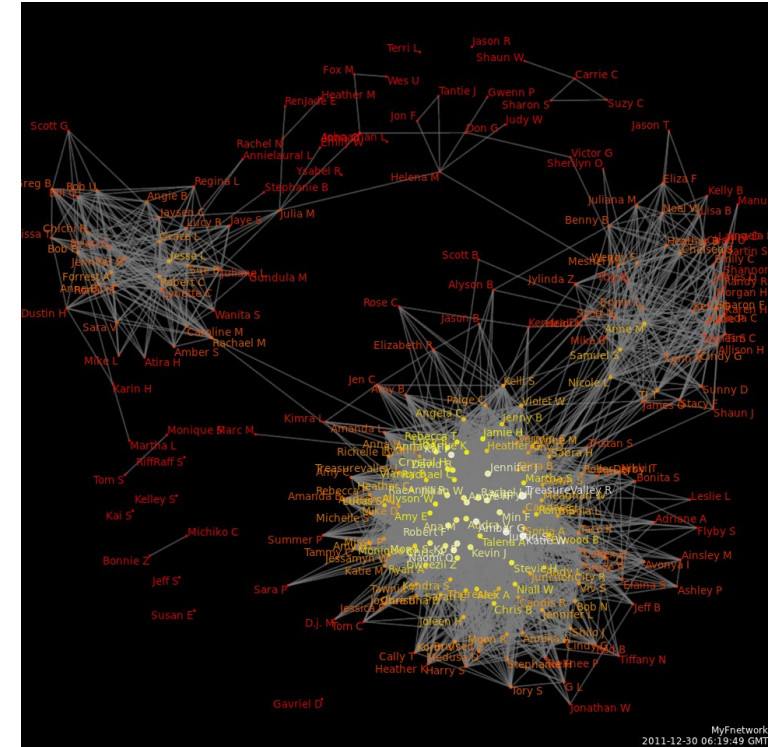

# Clustering

---

- Clustering is a classic unsupervised learning task.
  - .Organizing data
- There are many algorithms for clustering high-dimensional data

# Clustering

By Kencf0618 [CC BY-SA 3.0 (<https://creativecommons.org/licenses/by-sa/3.0/>) or CC BY-SA 3.0 (<https://creativecommons.org/licenses/by-sa/3.0/>)], from Wikimedia Commons



[https://en.wikipedia.org/wiki/Market\\_segmentation](https://en.wikipedia.org/wiki/Market_segmentation)


# Document Clustering

IBM

IBM Knowledge Center

[Content Classification](#) > [Content Classification 8.8.0](#) > [Configuring](#) > [Cat](#)  
Using the Taxonomy Proposer to discover new categories

Using the Taxonomy Proposer to discover new categories



- ❑ Data mining

- ❑ Often have huge numbers of documents

- ❑ How can we organize this?


- ❑ Key idea: documents are often in clusters

- ❑ Can we detect these clusters?

- ❑ Can be a lucrative service
  - See IBM service to left

# Clustering

- ❑ Clustering has many applications
  - Any time you want to segment data
  - Uncovering latent discrete variables
- ❑ Examples:
  - Segmenting sections of an image
  - Segmenting customers in market data



From: Market segmentation possibilities in the tourism market context of South Africa

# Image Segmentation

---



- ❑ Also from Bishop.
- ❑ Use K-means on the RGB values (dimension = 3)



# How can we find clusters in the data?

---

# What makes a “good” cluster?

$$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$$

$$c^{(1)}, c^{(2)}, \dots, c^{(N)} \quad 1 \leq c^{(i)} \leq K$$





Pair share: which clustering do you like better? Why?

Mathematically what makes one clustering assignment better than another?

# “Goodness” Metric

$$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$$

$$c^{(1)}, c^{(2)}, \dots, c^{(N)} \quad 1 \leq c^{(i)} \leq K$$



$$\sum_{\mathbf{x} \in \text{cluster 1}} \|\mathbf{x} - \mu_1\|_2^2 + \sum_{\mathbf{x} \in \text{cluster 2}} \|\mathbf{x} - \mu_2\|_2^2 + \sum_{\mathbf{x} \in \text{cluster 3}} \|\mathbf{x} - \mu_3\|_2^2$$

$$= \sum_{i=1}^3 \sum_{\mathbf{x} \in \text{cluster } i} \|\mathbf{x} - \mu_i\|_2^2$$


$$= \sum_{j=1}^N \|\mathbf{x}^{(j)} - \mu_{c^{(j)}}\|_2^2$$

Minimizes distortion  
function, J



# Goal: minimize our objective function

$$J(c, \mu) = \sum_{j=1}^N \|\mathbf{x}^{(j)} - \mu_{c(j)}\|_2^2 = \sum_{i=1}^2 \sum_{\mathbf{x} \in \text{cluster } i} \|\mathbf{x} - \mu_i\|_2^2$$



Pair share: Let  $K=2$ .

Where would you make the cluster centers:  $\mu_1, \mu_2$ ?

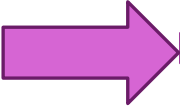
Pair share: For each point, which cluster would you assign it to?

Pair share What is  $J(c, \mu)$  for this cluster assignment?

$$J(c, \mu) = (1/2)^2 + (1/2)^2 + 2/9 + 2/9 + 5/9$$

# Outline

---

- ❑ Motivating Examples: Document clustering, image segmentation, image compression
- ❑ K-means
- ❑ K++-means (how to initialize the parameters before starting the algorithm)
- ❑ (On our own) K-means for document clustering

One clustering method is K-means clustering.

It finds a predetermined ( $K$ ) number of clusters in an unlabeled dataset

# K-Means


---

Assigns each example examples one of K clusters, where  $\mu_j$  is the center of cluster j (i.e., the *mean* of its cluster)  
 $c^{(i)}$  is the cluster  $x^{(i)}$  belongs to

$$J(c, \mu) = \sum_{i=1}^N \left\| x^{(i)} - \mu_{c^{(i)}} \right\|^2$$



NP hard to solve  
this problem!

There is an  
exponential number of  
ways to assign points  
to clusters



# Lloyd's Algorithm (Stuart Lloyd, 1957)

Pseudo code from CS229 Lecture notes



1. Initialize cluster *centroids*  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^d$  randomly
2. Repeat until convergence:  
For every  $i$ , set

$$c^{(i)} := \arg \min_j \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2$$

For every  $j \in \{1, \dots, K\}$ , set

$$\mu_j := \frac{\sum_{i=1}^N 1\{c^{(i)} = j\} \mathbf{x}^{(i)}}{\sum_{i=1}^N 1\{c^{(i)} = j\}}$$

}

Update cluster membership of every example. Every example belongs to the cluster it is closest to.



Suppose  $\mathbf{x}^{(2)}, \mathbf{x}^{(9)}, \mathbf{x}^{(21)}$  were assigned to cluster 1 then  $\mu_1 = (\mathbf{x}^{(2)} + \mathbf{x}^{(9)} + \mathbf{x}^{(21)})/3$

Definition:

$$1\{c^{(i)} = j\} = \begin{cases} 1 & c^{(i)} = j \\ 0 & c^{(i)} \neq j \end{cases}$$

# Lloyd's Algorithm (Stuart Lloyd, 1957)

Pseudo code from CS229 Lecture notes



1. Initialize cluster *centroids*  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^d$  randomly
2. Repeat until convergence:  
 For every  $i$ , set

$$c^{(i)} := \arg \min_j \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2$$

For every  $j \in \{1, \dots, K\}$ , set

$$\mu_j := \frac{\sum_{i=1}^N 1\{c^{(i)} = j\} \mathbf{x}^{(i)}}{\sum_{i=1}^N 1\{c^{(i)} = j\}}$$









}

Update cluster membership of every example. Every example belongs to the cluster it is closest to.

Update *centroid* of each cluster to be the *average(mean)* of examples assigned to cluster  $j$

Definition:

$$1\{c^{(i)} = j\} = \begin{cases} 1 & c^{(i)} = j \\ 0 & c^{(i)} \neq j \end{cases}$$



1. Initialize cluster *centroids*  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^d$  randomly
2. Repeat until convergence:
  - For every  $i$ , set

$$c^{(i)} := \arg \min_j \left\| x^{(i)} - \mu_j \right\|^2$$

For every  $j \in \{1, \dots, K\}$ , set

$$\mu_j := \frac{\sum_{i=1}^N 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^N 1\{c^{(i)} = j\}}$$

}

Definition:

$$1\{c^{(i)} = j\} = \begin{cases} 1 & c^{(i)} = j \\ 0 & c^{(i)} \neq j \end{cases}$$

# Centroid is Minimizer $\mu_j = \frac{1}{|S_j|} \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)}$

We show that our choice of  $\mu_j$  is better than any other point  $\mathbf{p}$ .

To show this we need to prove that:

$$\sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \frac{1}{|S_j|} \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)} \right\|^2 \leq \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mathbf{p} \right\|^2$$

Proof:

$$\sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mathbf{p} \right\|^2 = \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \underbrace{\mathbf{x}^{(i)} - \mu_j}_{\mathbf{a}} + \underbrace{\mu_j - \mathbf{p}}_{\mathbf{b}} \right\|^2$$

Adding  $0 = -\mu_j + \mu_j$

Here  $\mathbf{a}, \mathbf{b}$  are vectors. Notice that:  
 $\|\mathbf{a} + \mathbf{b}\|^2 = \|\mathbf{a}\|^2 + \|\mathbf{b}\|^2 + 2\mathbf{a}^T \mathbf{b}$

$$= \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2 + \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mu_j - \mathbf{p} \right\|^2 + 2 \sum_{\mathbf{x}^{(i)} \in S_j} (\mathbf{x}^{(i)} - \mu_j)^T (\mu_j - \mathbf{p})$$

$$= \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2 + \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mu_j - \mathbf{p} \right\|^2$$

$$\geq \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2$$

We can move  $(\mu_j - \mathbf{p})$  in front of the sum:  
 $2(\mu_j - \mathbf{p})^T \sum_{\mathbf{x}^{(i)} \in S_j} (\mathbf{x}^{(i)} - \mu_j)$

We can rewrite this as:

$$= 2(\mu_j - \mathbf{p})^T \left( \left( \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)} \right) - |S_j| \mu_j \right)$$

Now notice that:  $|S_j| \mu_j = \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)}$

Thus  $\left( \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)} \right) - |S_j| \mu_j = 0$



# Centroid is Minimizer $\mu_j = \frac{1}{|S_j|} \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)}$

We show that our choice of  $\mu_j$  is better than any other point  $\mathbf{p}$ .

To show this we need to prove that:

$$\sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2 \leq \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mathbf{p} \right\|^2$$

Proof:

$$\sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mathbf{p} \right\|^2 = \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \underbrace{\mathbf{x}^{(i)} - \mu_j}_{\mathbf{a}} + \underbrace{\mu_j - \mathbf{p}}_{\mathbf{b}} \right\|^2$$

Adding  $0 = -\mu_j + \mu_j$

Here  $\mathbf{a}, \mathbf{b}$  are vectors. Notice that:  
 $\|\mathbf{a} + \mathbf{b}\|^2 = \|\mathbf{a}\|^2 + \|\mathbf{b}\|^2 + 2\mathbf{a}^T \mathbf{b}$

$$= \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2 + \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mu_j - \mathbf{p} \right\|^2 + 2 \sum_{\mathbf{x}^{(i)} \in S_j} (\mathbf{x}^{(i)} - \mu_j)^T (\mu_j - \mathbf{p})$$

$$= \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2 + \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mu_j - \mathbf{p} \right\|^2$$

$$\geq \sum_{\mathbf{x}^{(i)} \in S_j} \left\| \mathbf{x}^{(i)} - \mu_j \right\|^2$$

We can move  $(\mu_j - \mathbf{p})$  in front of the sum:  
 $2(\mu_j - \mathbf{p})^T \sum_{\mathbf{x}^{(i)} \in S_j} (\mathbf{x}^{(i)} - \mu_j)$

We can rewrite this as:


$$= 2(\mu_j - \mathbf{p})^T \left( \left( \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)} \right) - |S_j| \mu_j \right)$$

Now notice that:  $|S_j| \mu_j = \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)}$

Thus  $\left( \sum_{\mathbf{x}^{(i)} \in S_j} \mathbf{x}^{(i)} \right) - |S_j| \mu_j = 0$


# K-Means illustrated

By Chire [GFDL (<http://www.gnu.org/copyleft/fdl.html>) or CC BY-SA 4.0 (<https://creativecommons.org/licenses/by-sa/4.0>)], from Wikimedia Commons



# Uh Oh...

- ❑ The K-means clustering algorithm is guaranteed to improve the result on each step...and converge - but not to a globally optimal solution.
- ❑ However, K-means is not guaranteed to find a global minimum - only a local minimum.
- ❑ Finding the global minimum K-means error is NP-hard...
- ❑ Run the algorithm with many initial configurations and keep the one that performs best



# E-M Algorithm

- ❑ The K-means algorithm is a variant of the E-M algorithm.
  - The E step (Expectation step) involves updating our expectation of what cluster each example belongs to.
  - The M step (Maximization step) involves maximizing the best location of the cluster centers.
- ❑ The algorithm works by minimizing a complex error function by separating the data into two steps: If one step is known, it is easy to optimize the other step

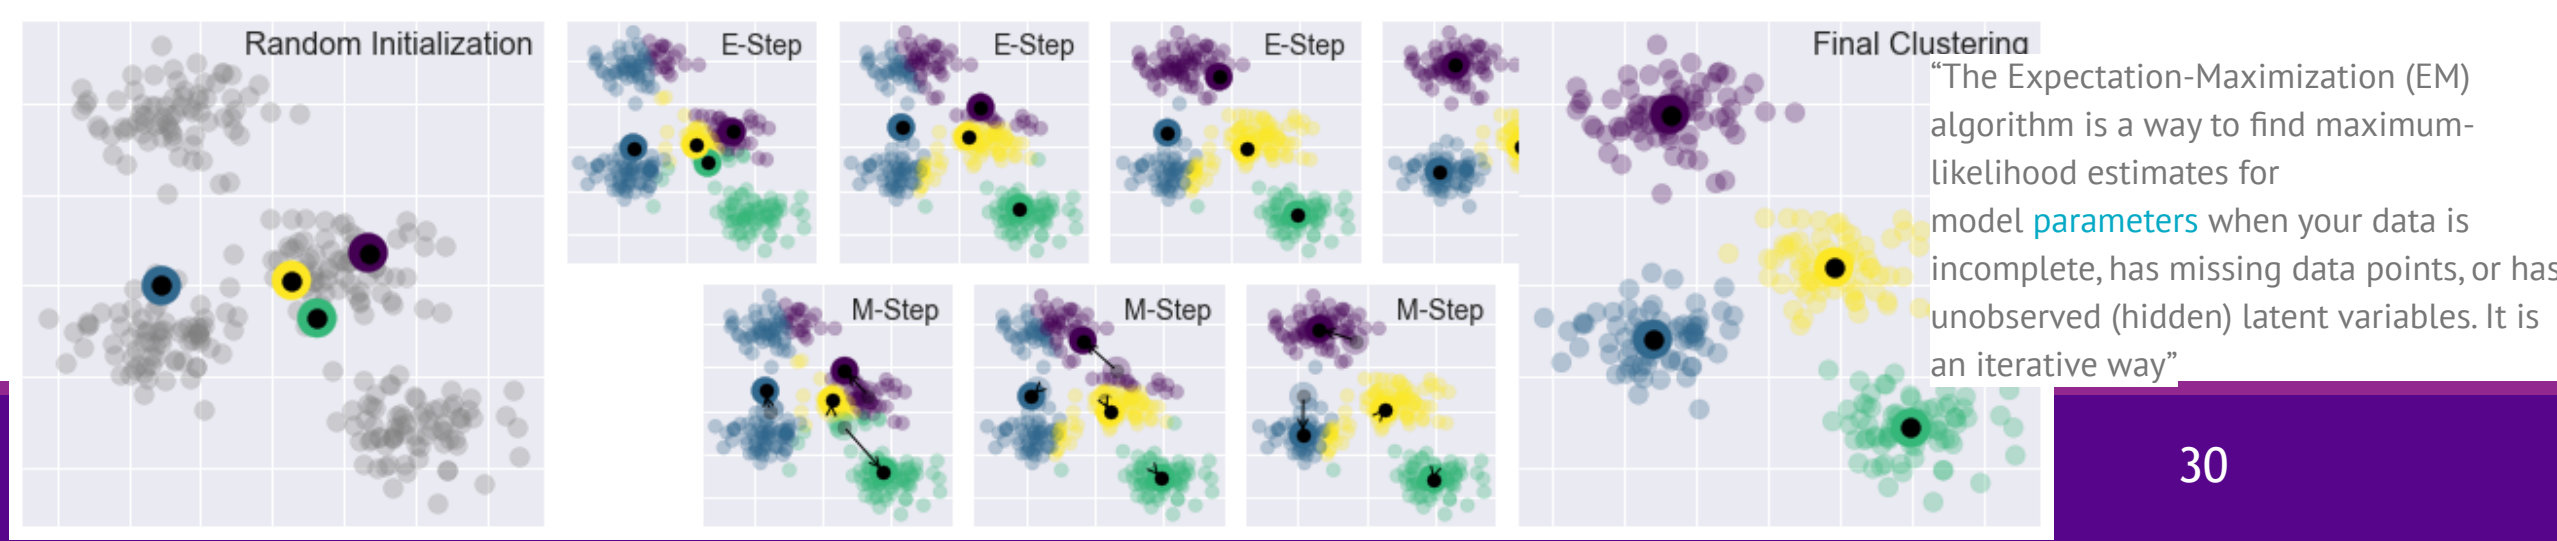
Given an assignment of points to clusters, could we find a better cluster center than taking the average of the points in a cluster to be its center?

- A) Yes
- B) No
- C) Maybe

# K-Means Converges

□ The algorithm converges to a partition that is “locally optimal.”

- Given the cluster centers  $\mu_j$ , we cannot find a better assignment of the examples to clusters.
- Given the cluster assignments ( $c^{(i)}$  for all  $i \in 1 \dots N$ ), we cannot find better centers.



# Proof of convergence (to a local min)

## Theorem (K - Means Convergence Theorem)

We update  $\mu$  and we update  $c$ . For each update we show that they never increase the value of

$$J(c, \mu) = \sum_{i=1}^N \left\| \mathbf{x}^{(i)} - \mu_{c^{(i)}} \right\|^2$$

There are only a finite number of values that can be assigned to  $\mu$  and  $c$ . ( $\mu$  is the mean of a subset of the examples and  $c \in \{1, 2, \dots, K\}$ ).

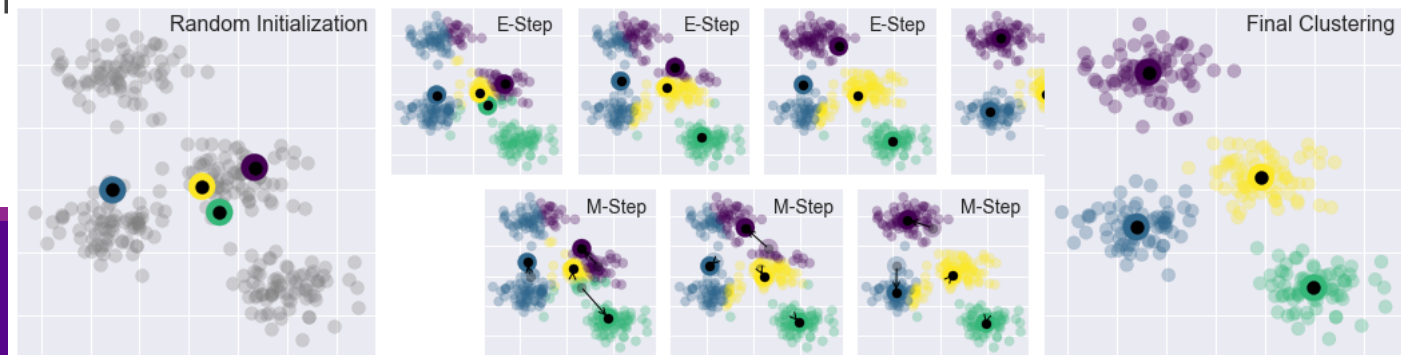
We also know that  $J(c, \mu) \geq 0$ .

Thus  $J(c, \mu)$  can only decrease a finite number of times. When it stops decreasing the algorithm has converged (to a local minimum)

When we update  $c^{(i)}$ , it must be that  $\left\| \mathbf{x}^{(i)} - \mu_{c^{(i \text{ new})}} \right\|^2 \leq \left\| \mathbf{x}^{(i)} - \mu_{c^{(i)}} \right\|^2$

When we update  $\mu_j$  as the mean of the points which are in this cluster - it directly minimizes  $\sum_{c^{(i)}=j} (\mathbf{x}^{(i)} - \mu_j)^2$

Thus every iteration decreases the cost function

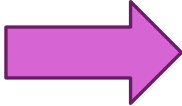


Since it is possible to converge to a local minimum instead of a global minimum, you should run the algorithm 10 times and choose the clustering with the lowest  $J(c, \mu)$



# Outline

---

- ❑ Motivating Examples: Document clustering, image segmentation, image compression
- ❑ K-means
-  ❑ K++-means (how to initialize the parameters before starting the algorithm)
- ❑ K-means for document clustering

The big concern is poor initialization at the start of the algorithm.



# How to choose the initial values...

□ One heuristic (we will refine it on the next slide) is to use the *furthest-first* algorithm

1. Pick a random example  $j$  and set  $\mu_1 = \mathbf{x}^{(j)}$

2. For  $k'' = 2..K$ :


Find the example  $j$  that is as far as possible from all previously selected means; namely:

$$j = \arg \max_j \min_{k' < k''} \left\| \mathbf{x}^{(j)} - \mu_{k'} \right\|^2$$


and set  $\mu_{k''} = \mathbf{x}^{(j)}$

Find index of the training example is farthest from its closest center


The problem is that this algorithm is sensitive to outliers.



# Outliers...



Instead of choosing the furthest example from your existing clusters, select the next center randomly with probability proportional to its distance squared.



Pair share: What are the distances?  
Compute one of the probabilities.

# K-means++ algorithm

## □ Algorithm k-means++

$\mu_1 = \mathbf{x}^{(j)}$  for  $j$  chosen uniformly at random // randomly initialize first point

for  $k''=2$  to  $K$  do

$$d_j = \min_{k' < k''} \|\mathbf{x}^{(j)} - \mu_{k'}\|, \forall j \quad // \text{compute distances}$$


$$p_j = \frac{d_j^2}{\sum_{i=1}^m d_i^2}, \forall j \quad // \text{normalize to probability distribution}$$

$j$  = random chosen with probability  $p_j$

$$\mu_{k''} = \mathbf{x}^{(j)}$$

Next run k-means using  $\mu$  as initial centers


After we find the initial centers - we run the K-means algorithm discussed earlier.



It can be proven that the expected value of the  $J(c, \mu)$  when running K-means++ is never more than  $O(\log K)$  times optimal  $J(c, \mu)$

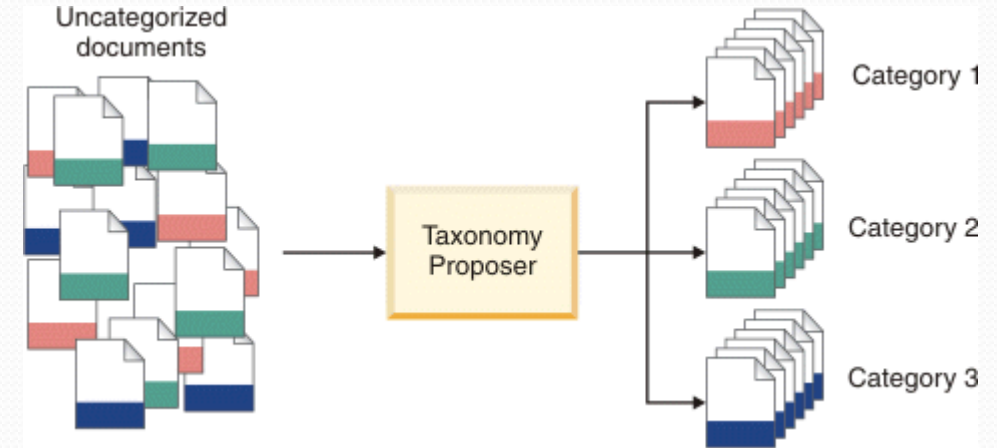
# Outline

---

- ❑ Motivating Examples: Document clustering, image segmentation, image compression
- ❑ K-means
- ❑ K++-means (how to initialize the parameters before starting the algorithm)
-  ❑ (On your own) K-means for document clustering



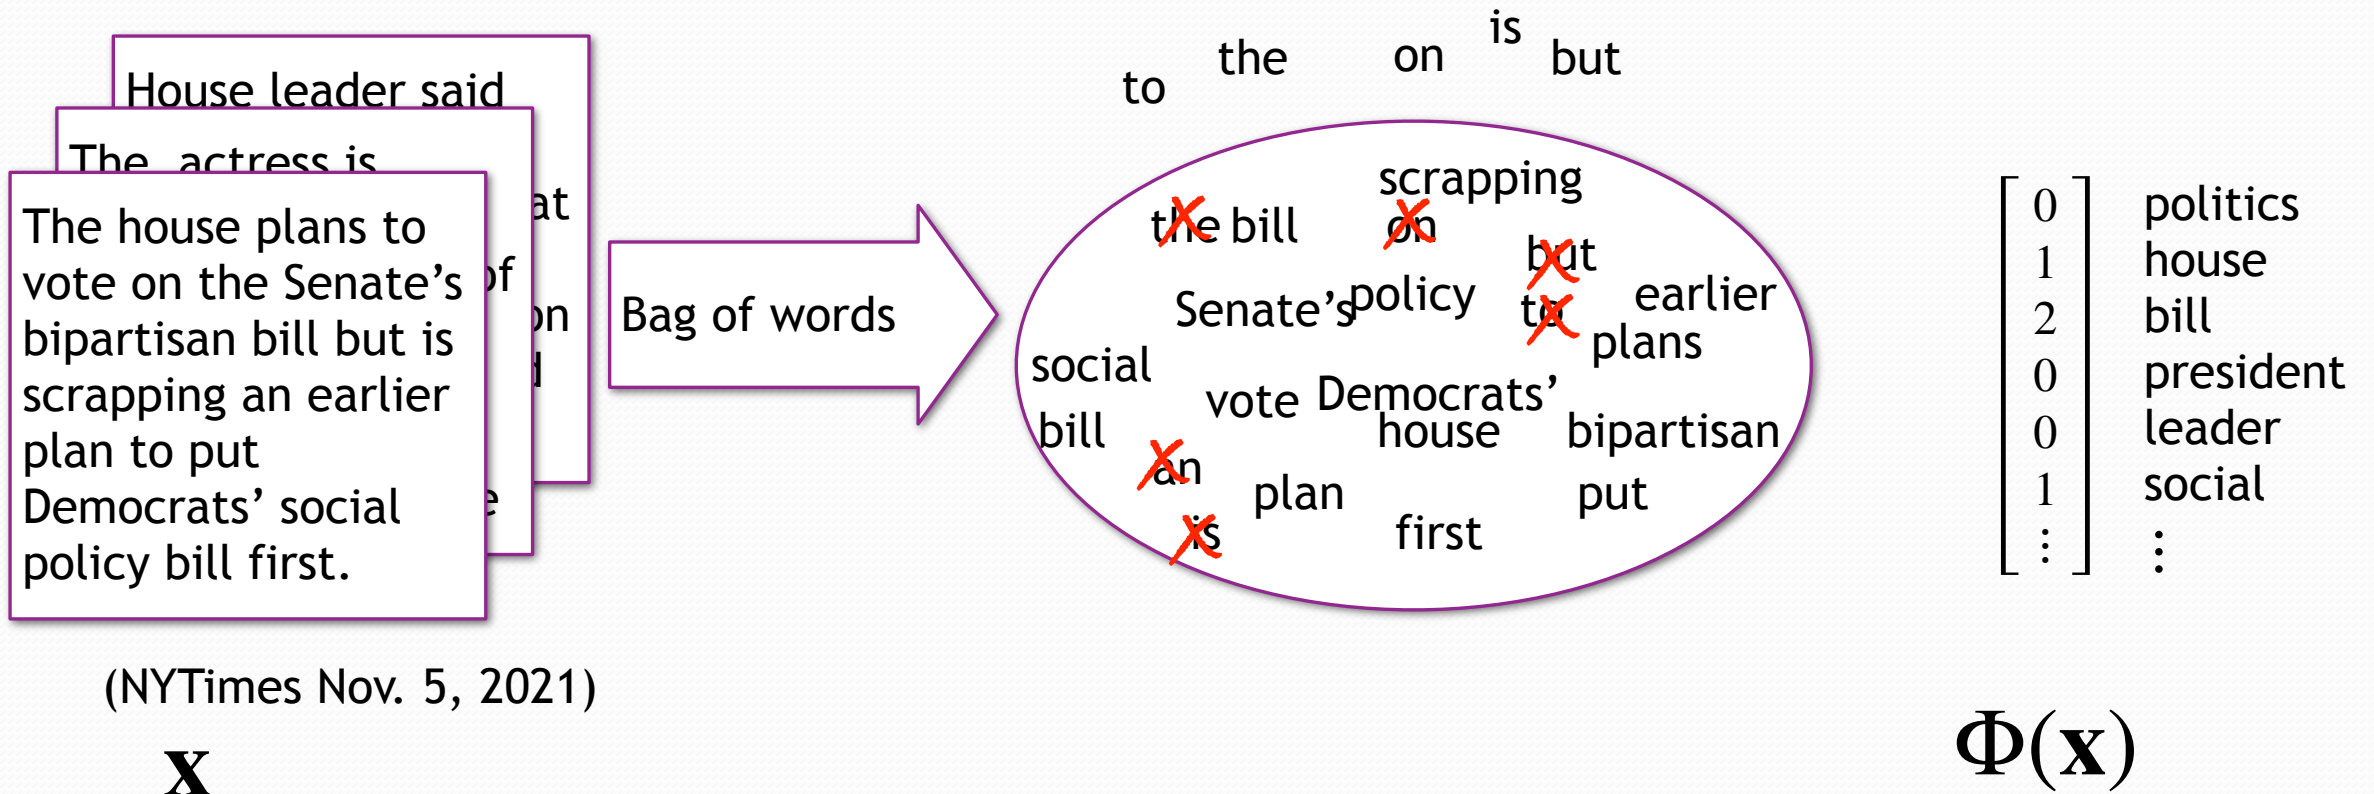
# Feature Extraction:



If we want to use K-means into cluster documents, we first need to convert text into a set of numerical values.

How can we do this?

# Documents as feature vectors



Transform the feature vectors to emphasize more “relevant” words

# Turning text into a feature vector

---

## Document 1

The quick brown  
fox jumped over  
the lazy dog's  
back.

## Document 2

Now is the time  
for all good men  
to come to the  
aid of their party.

- ❑ Document is natively text
- ❑ Must represent as a numeric vector
- ❑ Represent by word counts
  - Enumerate all words
  - Each document is count of frequencies
- ❑ Stopwords

# Discussion Questions

---

- ❑ Is the absolute number of times a word appears the correct metric?
- ❑ What about the length of the document?
- ❑ What about the frequency of the word?
- ❑ What words “matter”?

the, for, a, in

convolutional, gradient

- ❑ Perhaps:
  - if a word appears frequently, it is important (give it a high score)
  - If a word appears in many documents, it is not important (give it a low score)

# Ideas:

$$TF_{\text{"this"}, d_1} = \frac{1}{5} = 0.2$$

$$TF_{\text{"this"}, d_2} = \frac{1}{7} \approx 0.14$$

$$IDF_{\text{"this"}} = \log\left(\frac{2}{2}\right) = 0$$

$$TF_{\text{"example"}, d_1} = \frac{0}{5} = 0$$

$$TF_{\text{"example"}, d_2} = \frac{3}{7} \approx 0.429$$

$$IDF_{\text{"example"}} = \log\left(\frac{2}{1}\right) = 1$$

Example modified from <https://en.wikipedia.org/wiki/Tf%E2%80%93idf>

❑ How can we categorize how important a word is in a document?

❑ Perhaps:

- if a word appears frequently, it is important (give it a high score)
- except if the word appears in many documents, it is not important (give it a low score)

❑ Steps:

- Count the frequency of every word in the document

Term frequency

$$TF_{i,n} = \frac{\text{num times word } i \text{ in doc } n}{\text{total num words in doc } n}$$

- Determine how much information a word provides: Inverse Document Frequency (IDF)

The more common a word is  
the lower its IDF score

Inverse doc frequency

$$IDF_i = \log\left[\frac{\text{Total num docs in corpus}}{\text{Num docs with word } i}\right]$$

Document 1

Term	Term Count
this	1
Is	1
a	2
sample	1

Document 2

Term	Term Count
this	1
Is	1
another	2
example	3

# Ideas:

$$TF_{\text{"this"}, d_1} = \frac{1}{5} = 0.2$$

$$TF_{\text{"this"}, d_2} = \frac{1}{7} \approx 0.14$$

$$IDF_{\text{"this"}} = \log\left(\frac{2}{2}\right) = 0$$

$$TF_{\text{"example"}, d_1} = \frac{0}{5} = 0$$

$$TF_{\text{"example"}, d_2} = \frac{3}{7} \approx 0.429$$

$$IDF_{\text{"example"}} = \log\left(\frac{2}{1}\right) = 1$$

Frequency is relative to the size of the document

Example modified from

<https://www.wikipedia.org/wiki/Tf%E2%80%93idf>

How can we categorize how important a word is in a document?

Perhaps:

- if a word appears frequently, it is important (give it a high score)
- except if the word appears in many documents, it is not important (give it a low score)

Steps:

- Count the frequency of every word in the document

Term frequency

$$TF_{i,n} = \frac{\text{num times word } i \text{ in doc } n}{\text{total num words in doc } n}$$

- Determine how much information a word provides: Inverse Document Frequency (IDF)

The more common a word is the lower its IDF score

Inverse doc frequency

$$IDF_i = \log\left[\frac{\text{Total num docs in corpus}}{\text{Num docs with word } i}\right]$$

Document 1

Term	Term Count
this	1
Is	1
a	2
sample	1

Document 2

Term	Term Count
this	1
Is	1
another	2
example	3

# Term Frequency - Inverse Document Frequency

Use TF-IDF weight for vectors:

$$X[n, i] = TF_{i,n} \times IDF_i$$

Document weight vector

Term frequency

Inverse doc frequency

$$TF_{i,n} = \frac{\text{num times word } i \text{ in doc } n}{\text{total num words in doc } n}$$

$$IDF_i = \log \left[ \frac{\text{Total num docs in corpus}}{\text{Num docs with word } i} \right]$$

$$TF_{\text{"this"}, d_1} = \frac{1}{5} = 0.2$$

$$IDF_{\text{"this"}} = \log \left( \frac{2}{2} \right) = 0$$

$$TF_{\text{"example"}, d_1} = \frac{0}{5} = 0$$

$$IDF_{\text{"example"}} = \log \left( \frac{2}{1} \right) = 1$$

$$TF_{\text{"this"}, d_2} = \frac{1}{7} \approx 0.14$$

$$TF_{\text{"example"}, d_2} = \frac{3}{7} \approx 0.429$$

Example modified from <https://en.wikipedia.org/wiki/Tf%E2%80%93idf>

$$TF\text{-}IDF_{\text{"this"}, d_1} = 0.2 \times 0 = 0$$

$$TF\text{-}IDF_{\text{"example"}, d_1, D} = 0. \times 1 = 0$$

$$TF\text{-}IDF_{\text{"this"}, d_2} = 0.14 \times 0 = 0$$

$$TF\text{-}IDF_{\text{"example"}, d_2} = 0.429 \times 1 = 0.429$$